

CSC 4304 - Systems Programming
Fall 2010

LECTURE - VII
UNIX PROCESS ENVIRONMENT

Tevfik Koşar

Louisiana State University
September 14th, 2010

In Today's Class

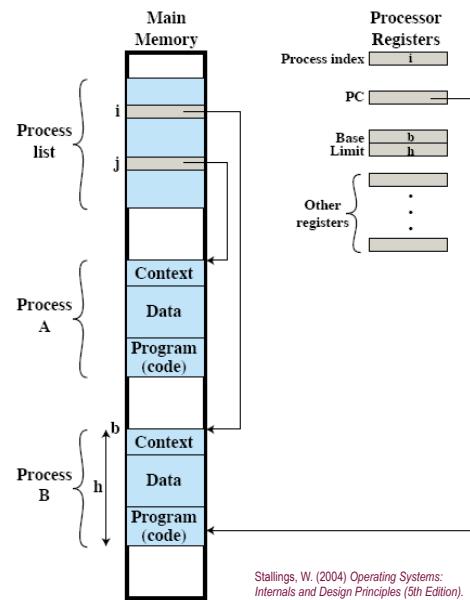
- Unix Process Environment
 - Process Concept
 - Creation & Termination of Processes
 - Exec() & Fork()
 - ps -- get process info
 - Shell & its implementation

Process Concept

- a **Process** is a program in execution;

➤ A process image consists of three components

- user address space
1. an executable program
 2. the associated data needed by the program
 3. the execution context of the process, which contains all information the O/S needs to manage the process (ID, state, CPU registers, stack, etc.)



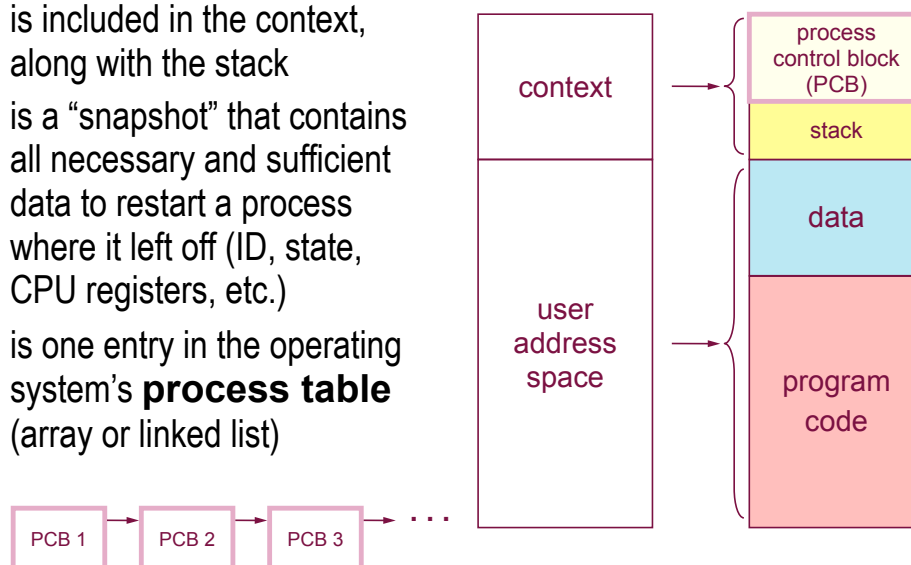
Typical process image implementation

Process Control Block

➤ The Process Control Block (PCB)

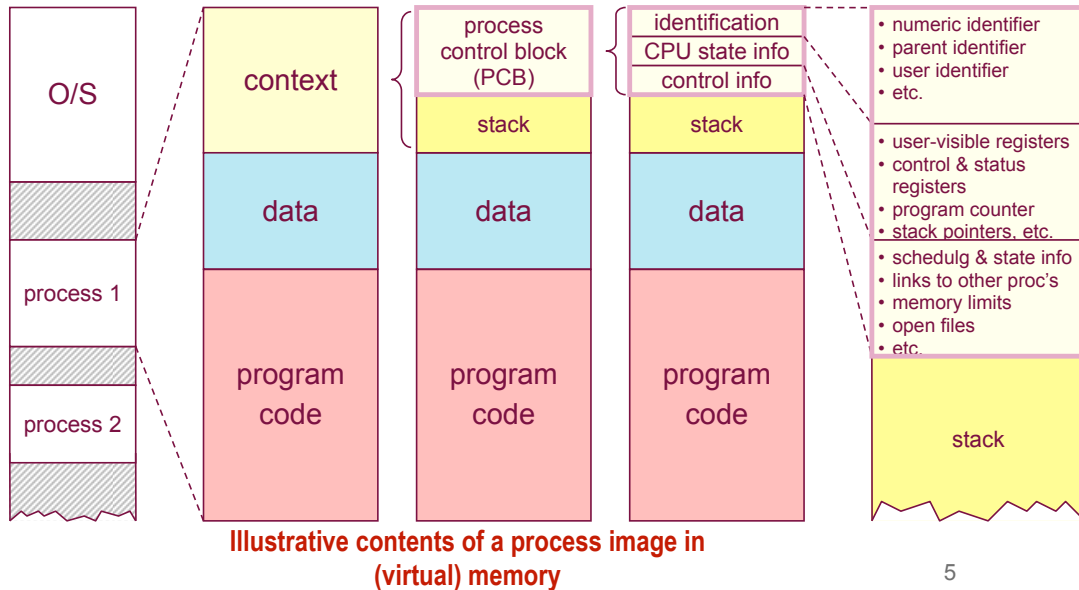
- ✓ is included in the context, along with the stack
- ✓ is a “snapshot” that contains all necessary and sufficient data to restart a process where it left off (ID, state, CPU registers, etc.)
- ✓ is one entry in the operating system’s **process table** (array or linked list)

Typical process image implementation




Process Control Block

➤ Example of process and PCB location in memory



Process State

- As a process executes, it changes *state*
 - new:** The process is being created
 - ready:** The process is waiting to be assigned to a process
 - running:** Instructions are being executed
 - waiting:** The process is waiting for some event to occur
 - terminated:** The process has finished execution



\$ ps

PID	TTY	TIME	CMD
18684	pts/4	00:00:00	bash
18705	pts/4	00:00:00	ps

\$ ps a

PID	TTY	STAT	TIME	COMMAND
6702	tty7	Ss+	15:10	/usr/X11R6/bin/X :0 -audit 0
7024	tty1	Ss+	0:00	/sbin/mingetty --noclear tty1
7025	tty2	Ss+	0:00	/sbin/mingetty tty2
7026	tty3	Ss+	0:00	/sbin/mingetty tty3
7027	tty4	Ss+	0:00	/sbin/mingetty tty4
7028	tty5	Ss+	0:00	/sbin/mingetty tty5
7029	tty6	Ss+	0:00	/sbin/mingetty tty6
17166	pts/6	Ss	0:00	-bash
17191	pts/6	S+	0:00	pico program3.cc
17484	pts/5	Ss+	0:00	-bash
17555	pts/7	Ss+	0:00	-bash
17646	pts/8	Ss	0:00	-bash
17809	pts/10	Ss	0:00	-bash
17962	pts/8	S+	0:00	pico prog2.java
17977	pts/1	Ss	0:00	-bash
18014	pts/9	Ss+	0:00	-bash
18259	pts/10	T	0:00	a.out
18443	pts/2	Ss	0:00	-bash
18511	pts/1	S+	0:00	pico program3.cc
18684	pts/4	Ss	0:00	-bash

\$ ps la

F	UID	PID	PPID	PRI	NI	VSZ	RSS	WCHAN	STAT	TTY	TIME	COMMAND
4	0	6702	6701	15	0	25416	7204	-	Ss+	tty7	15:10	/usr/X11R6/bin/X :0 -
audit 0 -auth /var/lib/g												
4	0	7024	1	17	0	3008	4	-	Ss+	tty1	0:00	/sbin/mingetty --
noclear tty1												
4	0	7025	1	16	0	3008	4	-	Ss+	tty2	0:00	/sbin/mingetty tty2
4	0	7026	1	16	0	3012	4	-	Ss+	tty3	0:00	/sbin/mingetty tty3
4	0	7027	1	17	0	3008	4	-	Ss+	tty4	0:00	/sbin/mingetty tty4
4	0	7028	1	17	0	3008	4	-	Ss+	tty5	0:00	/sbin/mingetty tty5
4	0	7029	1	17	0	3008	4	-	Ss+	tty6	0:00	/sbin/mingetty tty6
0	2317	17166	17165	15	0	9916	2300	wait	Ss	pts/6	0:00	-bash
0	2317	17191	17166	16	0	8688	1264	-	S+	pts/6	0:00	pico program3.cc
0	2238	17484	17483	16	0	9916	2300	-	Ss+	pts/5	0:00	-bash
0	2611	17555	17554	15	0	9912	2292	-	Ss+	pts/7	0:00	-bash
0	2631	17646	17644	16	0	9912	2300	wait	Ss	pts/8	0:00	-bash
0	2211	17809	17808	15	0	9916	2324	wait	Ss	pts/10	0:00	-bash
0	2631	17962	17646	16	0	8688	1340	-	S+	pts/8	0:00	pico prog2.java
0	2320	17977	17976	16	0	9912	2304	wait	Ss	pts/1	0:00	-bash

\$ ps -ax

PID	TTY	STAT	TIME	COMMAND
1	?	S	0:02	init [5]
2	?	S	0:00	[migration/0]
3	?	SN	0:00	[ksoftirqd/0]
4	?	S	0:00	[migration/1]
5	?	SN	0:01	[ksoftirqd/1]
6	?	S	0:00	[migration/2]
7	?	SN	0:16	[ksoftirqd/2]
8	?	S	0:00	[migration/3]
9	?	SN	0:16	[ksoftirqd/3]
10	?	S<	0:00	[events/0]
11	?	S<	0:00	[events/1]
12	?	S<	0:00	[events/2]
13	?	S<	0:00	[events/3]
14	?	S<	0:00	[khelper]
15	?	S<	0:00	[kthread]
653	?	S<	0:00	[kacpid]
994	?	S<	0:00	[kblockd/0]
995	?	S<	0:00	[kblockd/1]
996	?	S<	0:01	[kblockd/2]
997	?	S<	0:00	[kblockd/3]

Process Creation

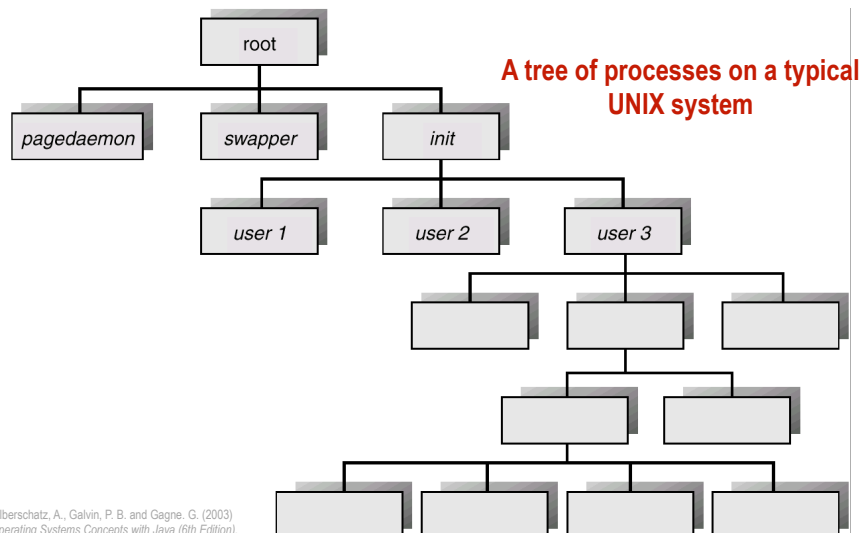
➤ Some events that lead to process creation (enter)

- all cases of process spawning
- ✓ the system boots
 - when a system is initialized, several background processes or “daemons” are started (email, logon, etc.)
 - ✓ a user requests to run an application
 - by typing a command in the CLI shell or double-clicking in the GUI shell, the user can launch a new process
 - ✓ an existing process spawns a child process
 - for example, a server process (print, file) may create a new process for each request it handles
 - the *init* daemon waits for user login and spawns a shell
 - ✓ a batch system takes on the next job in line

11

Process Creation

➤ Process creation by spawning



12

Process Creation

```
...
int main(...)
{
    ...
    if ((pid = fork()) == 0)                // create a process
    {
        fprintf(stdout, "Child pid: %i\n", getpid());
        err = execvp(command, arguments);    // execute child
                                              // process
        fprintf(stderr, "Child error: %i\n", errno);
        exit(err);
    }
    else if (pid > 0)                       // we are in the
    {                                       // parent process
        fprintf(stdout, "Parent pid: %i\n", getpid());
        pid2 = waitpid(pid, &status, 0);    // wait for child
        ...                                // process
    }
    ...
    return 0;
}
```

13


Process Creation

1. Clone child process

✓ pid = fork()

2. Replace child's image

✓ execve(name, ...)



14

Fork Example 1

```
#include <stdio.h>

main()
{
    int ret_from_fork, mypid;

    mypid = getpid();          /* who am i? */
    printf("Before: my pid is %d\n", mypid); /* tell pid */

    ret_from_fork = fork();

    sleep(1);
    printf("After: my fork returns pid : %d, said %d\n",
           ret_from_fork, getpid());
}
```

15

Fork Example 2

```
#include <stdio.h>

main()
{
    fork();
    fork();
    fork();
    printf("my pid is %d\n", getpid() );
}
```

How many lines of output will this produce?

16

Shell

- A tool for process and program control
- Three main functions
 - Shells run programs
 - Shells manage I/O
 - Shells can be programmed
- Main Loop of a Shell

```
while (!end_of_input){  
    get command  
    execute command  
    wait for command to finish  
}
```

17

How does a Program run another Program?

- Program calls **execvp**

```
int execvp(const char *file, char *const argv[]);
```

- Kernel loads program from disk into the process
- Kernel copies arglist into the process
- Kernel calls `main(argc,argv)`

18

Exec Family

```
int execl(const char *path, const char *arg, ...);

int execvp(const char *file, const char *arg, ...);

int execle(const char *path, const char *arg , ...,
           char * const envp[]);

int execv(const char *path, char *const argv[]);

int execvp(const char *file, char *const argv[]);
```

19

execvp is like a Brain Transplant

- **execvp loads the new program into the current process, replacing the code and data of that process!**

20

Running “ls -l”

```
#include <unistd.h>
#include <stdio.h>

main()
{
    char    *arglist[3];

    arglist[0] = "ls";
    arglist[1] = "-l";
    arglist[2] = 0 ;

    printf(" * * About to exec ls -l\n");
    execvp( "ls" , arglist );
    printf(" * * * ls is done. bye\n");
}
```

21

Writing a Shell v1.0

```
int main()
{
    char *arglist[MAXARGS+1];    /* an array of ptrs    */
    int numargs;                 /* index into array */
    char argbuf[ARGLEN];         /* read stuff here */
    char *makestring();          /* malloc etc      */

    numargs = 0;
    while ( numargs < MAXARGS )
    {
        printf("Arg[%d]? ", numargs);
        if ( fgets(argbuf, ARGLEN, stdin) && *argbuf != '\n' )
            arglist[numargs++] = makestring(argbuf);
        else
        {
            if ( numargs > 0 ){          /* any args? */
                arglist[numargs]=NULL; /* close list */
                execute( arglist );    /* do it */
                numargs = 0;           /* and reset */
            }
        }
    }
    return 0;
}
```

```
#include <stdio.h>
#include <signal.h>
#include <string.h>

#define MAXARGS 20
#define ARGLEN 100
```

22

Writing a Shell v1.0 (cont.)

```
int execute( char *arglist[] )
{
    execvp(arglist[0], arglist);          /* do it */
    perror("execvp failed");
    exit(1);
}

char * makestring( char *buf )
{
    char *cp, *malloc();

    buf[strlen(buf)-1] = '\0';            /* trim newline */
    cp = malloc( strlen(buf)+1 );         /* get memory */
    if ( cp == NULL ){                    /* or die */
        fprintf(stderr, "no memory\n");
        exit(1);
    }
    strcpy(cp, buf);                      /* copy chars */
    return cp;                            /* return ptr */
}
```

23

Writing a Shell v2.0

```
execute( char *arglist[] )
{
    int pid, exitstatus;                  /* of child */

    pid = fork();                         /* make new process */
    switch( pid ){
        case -1:
            perror("fork failed");
            exit(1);
        case 0:
            execvp(arglist[0], arglist);   /* do it */
            perror("execvp failed");
            exit(1);
        default:
            while( wait(&exitstatus) != pid )
                ;
            printf("child exited with status %d,%d\n",
                   exitstatus>>8, exitstatus&0377);
    }
}
```

24

Process Termination

➤ Some events that lead to process termination (exit)

- ✓ regular completion, with or without error code
 - process-triggered
 - the process voluntarily executes an **exit(err)** system call to indicate to the O/S that it has finished
- ✓ fatal error (uncatchable or uncaught)
 - O/S-triggered (following system call or preemption)
 - service errors: no memory left for allocation, I/O error, etc.
 - total time limit exceeded
 - hardware interrupt-triggered
 - arithmetic error, out-of-bounds memory access, etc.
- ✓ killed by another process via the kernel
 - software interrupt-triggered
 - the process receives a **SIGKILL** signal
 - in some systems the parent takes down its children with it

25

Exercise

Improve the Shell v2.0 by:

- Allow the user to type all the arguments on one line
- Allow the user to quit by typing exit

26

Summary

- Unix Process Environment
 - Process Concept
 - `ps` -- get process info
 - Shell & its implementation
 - `Exec()` & `Fork()`
 - Creation & Termination of Processes
- Next Class: Process Control
- Try “fork” and “shell” examples



27

Acknowledgments

- Advanced Programming in the Unix Environment by R. Stevens
- The C Programming Language by B. Kernighan and D. Ritchie
- Understanding Unix/Linux Programming by B. Molay
- Lecture notes from B. Molay (Harvard), T. Kuo (UT-Austin), G. Pierre (Vrije), M. Matthews (SC), and B. Knicki (WPI).

28